

Name : Shobhit Mirjankar

Batch : C32

Roll No : 1902101

```
In [17]: # Importing libraries

import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import classification_report
import warnings
warnings.filterwarnings('ignore')
```

```
In [18]: dataset = pd.read_csv('churn.csv')[0:201]
dataset
```

```
Out[18]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	France	Female	42	2	0.00	1	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	1	
2	502	France	Female	42	8	159660.80	3	1	0	
3	699	France	Female	39	1	0.00	2	0	0	
4	850	Spain	Female	43	2	125510.82	1	1	1	
...
196	616	Spain	Female	32	6	0.00	2	1	1	
197	721	Germany	Male	37	3	107720.64	1	1	1	
198	501	France	Male	57	10	0.00	2	1	1	
199	521	France	Male	35	6	96423.84	1	1	0	
200	850	Spain	Male	30	2	141040.01	1	1	1	

201 rows × 11 columns

```
In [19]: # EDA

print('Columns : ',list(dataset.columns))
print()
print('Number of missing values : ',dataset.isnull().sum().sum())
print()
dataset.describe()
```

```
Columns :  ['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']
```

Number of missing values : 0

Out[19]:	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000	201.000000
mean	640.467662	37.965174	5.174129	74504.686269	1.542289	0.676617	0.482587	99189.93
std	108.463313	9.763902	2.987395	62726.490142	0.591145	0.468935	0.500944	57641.05
min	376.000000	19.000000	0.000000	0.000000	1.000000	0.000000	0.000000	600.36
25%	553.000000	32.000000	2.000000	0.000000	1.000000	0.000000	0.000000	47125.11
50%	646.000000	36.000000	5.000000	96645.540000	2.000000	1.000000	0.000000	99449.86
75%	722.000000	43.000000	8.000000	125851.930000	2.000000	1.000000	1.000000	147132.46
max	850.000000	75.000000	10.000000	213146.200000	4.000000	1.000000	1.000000	199725.39

```
In [20]: # Data Analysis

print(dataset['Geography'].unique())
print(dataset['Gender'].unique())

['France' 'Spain' 'Germany']
['Female' 'Male']
```

```
In [21]: # Data Preprocessing

le = LabelEncoder()
dataset['Gender'] = le.fit_transform(dataset['Gender'])

dataset = pd.get_dummies(dataset, columns = ['Geography'])
dataset
```

Out[21]:	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exit
0	619	0	42	2	0.00	1	1	1	101348.88	
1	608	0	41	1	83807.86	1	0	1	112542.58	
2	502	0	42	8	159660.80	3	1	0	113931.57	
3	699	0	39	1	0.00	2	0	0	93826.63	
4	850	0	43	2	125510.82	1	1	1	79084.10	
...
196	616	0	32	6	0.00	2	1	1	43001.46	
197	721	1	37	3	107720.64	1	1	1	158591.12	
198	501	1	57	10	0.00	2	1	1	47847.19	
199	521	1	35	6	96423.84	1	1	0	10488.44	
200	850	1	30	2	141040.01	1	1	1	5978.20	

201 rows × 13 columns

```
In [23]: dataset[['Geography_France', 'Geography_Germany', 'Geography_Spain', 'EstimatedSalary', 'Exit']]
```

Out[23]:

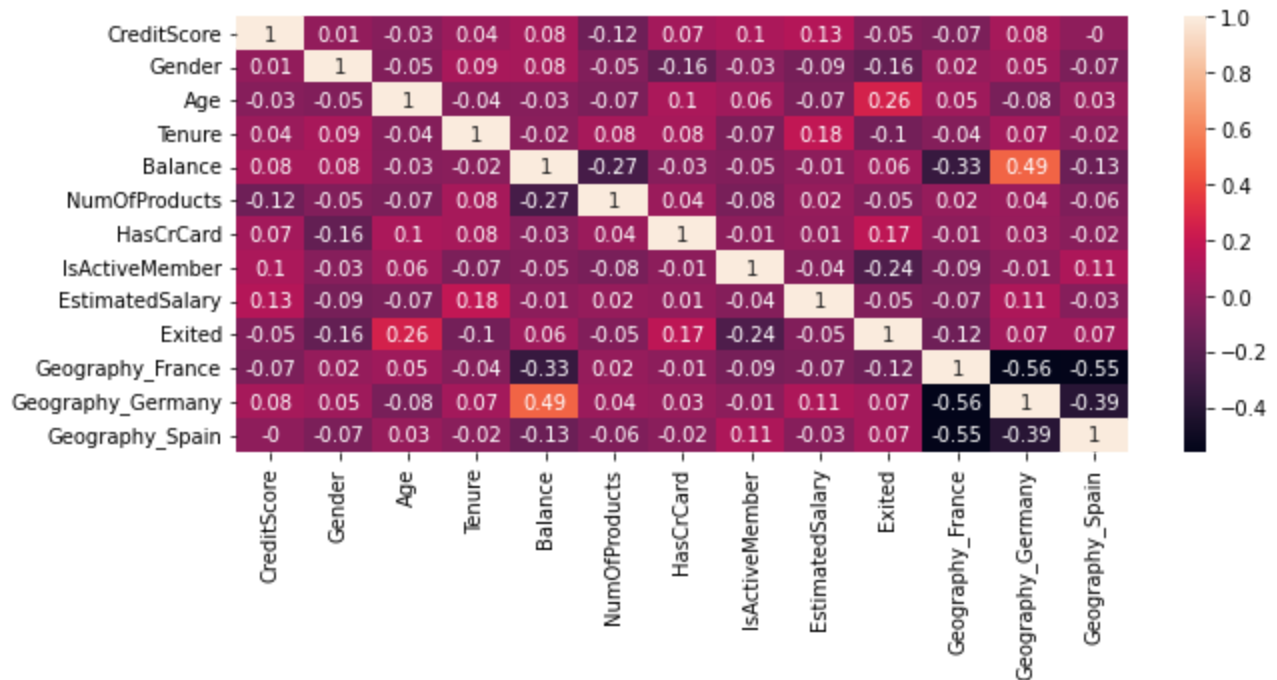
	Geography_France	Geography_Germany	Geography_Spain	EstimatedSalary	Exited
0	1	0	0	101348.88	1
1	0	0	1	112542.58	0
2	1	0	0	113931.57	1
3	1	0	0	93826.63	0
4	0	0	1	79084.10	0
...
196	0	0	1	43001.46	0
197	0	1	0	158591.12	0
198	1	0	0	47847.19	0
199	1	0	0	10488.44	0
200	0	0	1	5978.20	0

201 rows × 5 columns

In [6]:

```
# Feature Importance

plt.figure(figsize=(10, 4))
correl_matrix = dataset.corr().round(2)
sns.heatmap(data=correl_matrix, annot=True)
plt.show()
```



Hence we can see that most important features for predicting class label 'Exited' are : Age, IsActiveMember, HasCrCard, Gender, Geography_France

In [7]:

```
X = dataset[['Age', 'Geography_France', 'IsActiveMember', 'HasCrCard', 'Gender']]
y = dataset['Exited']

X
```

Out[7]:

	Age	Geography_France	IsActiveMember	HasCrCard	Gender
0	42		1	1	0
1	41		0	1	0
2	42		1	0	1
3	39		1	0	0
4	43		0	1	1
...
196	32		0	1	1
197	37		0	1	1
198	57		1	1	1
199	35		1	0	1
200	30		0	1	1

201 rows × 5 columns

In [8]:

```
y
```

Out[8]:

```
0      1
1      0
2      1
3      0
4      0
..
196    0
197    0
198    0
199    0
200    0
Name: Exited, Length: 201, dtype: int64
```

In [9]:

```
# Train test split

X_train1, X_test1, y_train1, y_test1 = train_test_split(X,y,test_size=0.1,
                                                         random_state=13)
X_train2, X_test2, y_train2, y_test2 = train_test_split(X,y,test_size=0.2,
                                                         random_state=13)
X_train3, X_test3, y_train3, y_test3 = train_test_split(X,y,test_size=0.3,
                                                         random_state=13)

X_train1
```

Out[9]:

	Age	Geography_France	IsActiveMember	HasCrCard	Gender
143	52		0	0	1
68	35		0	1	0
7	29		0	0	1
125	42		1	0	1
175	35		0	1	0
...

	Age	Geography_France	IsActiveMember	HasCrCard	Gender
98	22		0	0	1
16	58		0	1	1
74	36		1	0	1
176	30		1	1	0
82	36		0	0	0

180 rows × 5 columns

In [10]: y_train1

Out[10]:

```

143    1
68     0
7      1
125    1
175    0
..
98     0
16     1
74     0
176    0
82     0
Name: Exited, Length: 180, dtype: int64

```

Decision Tree Classifier (90-10 Split)

In [11]:

```

# Decision Tree Classifier for 90-10 Split

dtree = DecisionTreeClassifier(random_state=13)
dtree.fit(X_train1,y_train1)
y_pred_train1 = dtree.predict(X_train1)
y_pred_test1 = dtree.predict(X_test1)

print('90-10 Model performance on Training Set : \n')
print(classification_report(y_train1,y_pred_train1))
print()
print()
print('90-10 Model performance on Test Set : \n')
print(classification_report(y_test1,y_pred_test1))
print()
print()
X_ip = list(map(int,
    input("Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : ")
    .split()))[:5]
print('Predicted class : ',dtree.predict([X_ip])[0])
print()
print()
print('Result : ')
pd.DataFrame({'Actual':y_test1,'Predicted':y_pred_test1})

```

90-10 Model performance on Training Set :

	precision	recall	f1-score	support
0	0.96	0.99	0.98	147
1	0.96	0.82	0.89	33
accuracy			0.96	180

macro avg	0.96	0.91	0.93	180
weighted avg	0.96	0.96	0.96	180

90-10 Model performance on Test Set :

	precision	recall	f1-score	support
0	0.64	0.69	0.67	13
1	0.43	0.38	0.40	8
accuracy			0.57	21
macro avg	0.54	0.53	0.53	21
weighted avg	0.56	0.57	0.57	21

Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : 40 0 1 1 0
 Predicted class : 1

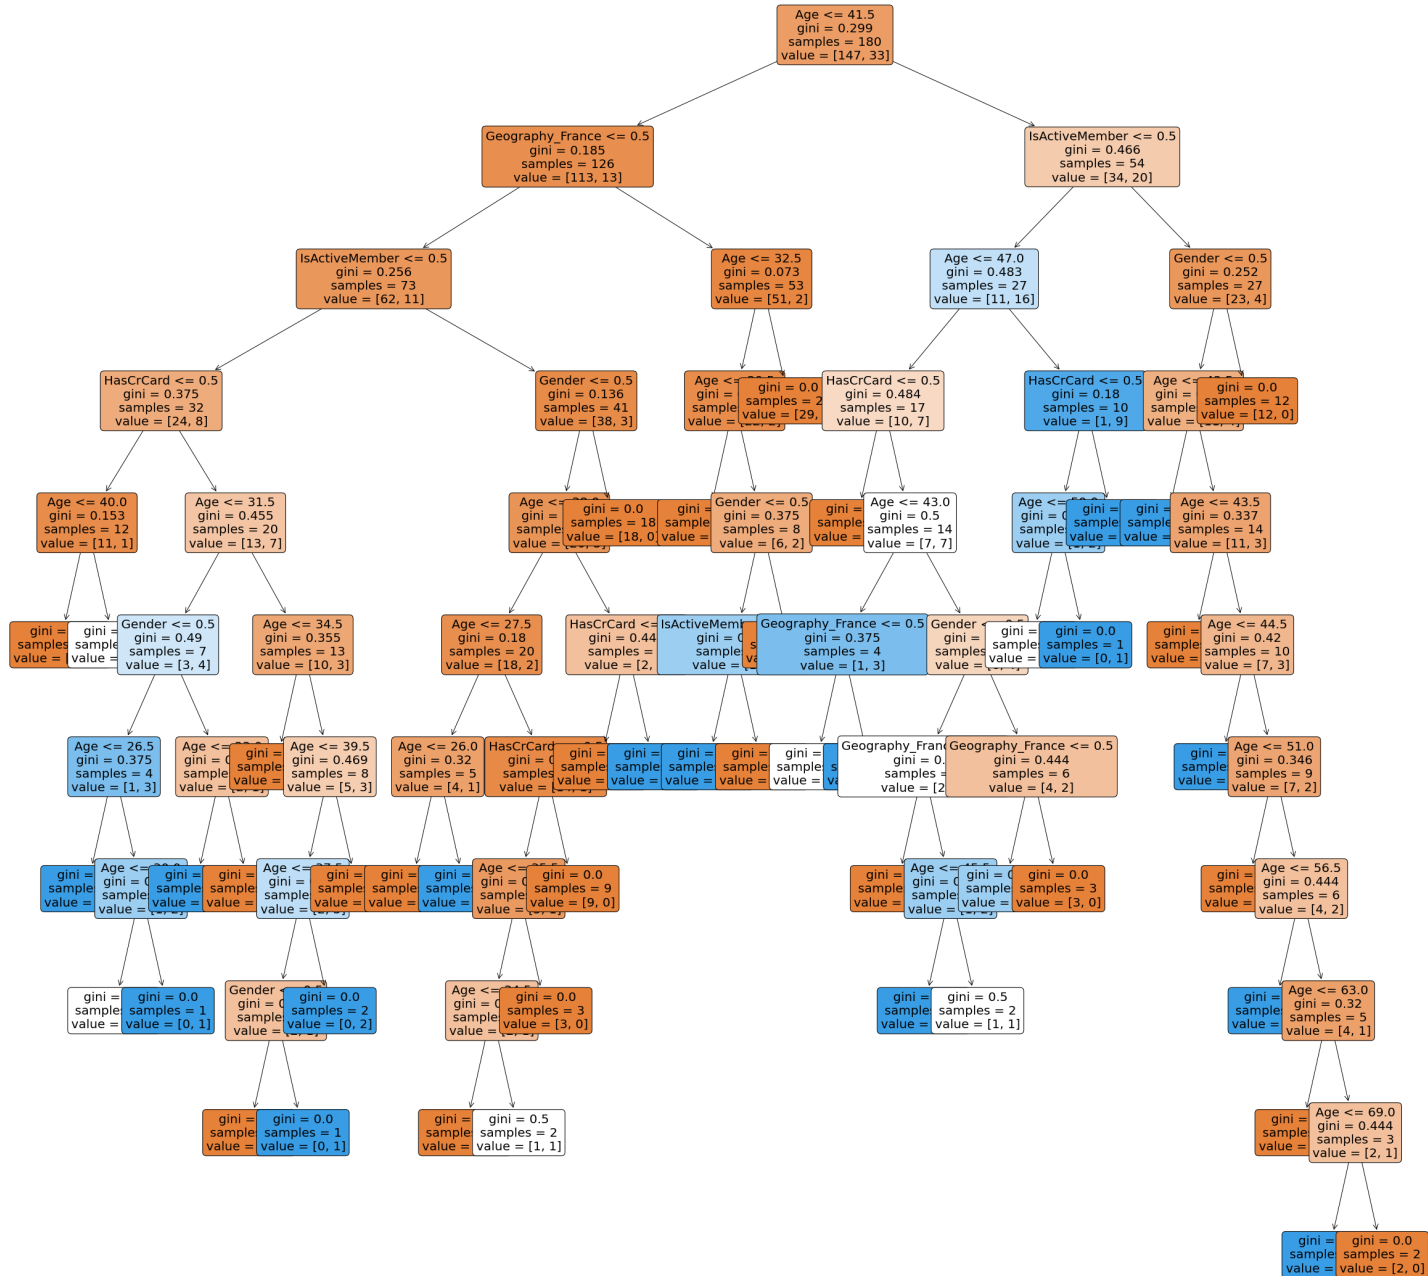
Result :

Out[11]:

	Actual	Predicted
101	0	0
87	0	0
23	0	0
140	0	1
114	1	0
190	1	0
65	0	0
111	0	1
33	0	1
196	0	0
194	0	0
30	1	1
70	1	1
169	0	0
164	1	0
132	0	0
103	0	1
179	1	1
105	1	0
184	1	0
191	0	0

In [12]: *# Visualize the Tree*

```
plt.figure(figsize=(40,40))
plot_tree(dtree, feature_names = X_train1.columns, filled=True, fontsize=20,
          rounded = True)
plt.show()
```



Training set gave 96% accuracy and test set gave 57% accuracy hence model is overfitted

Decision Tree Classifier (80-20 Split)

In [13]:

```
# Decision Tree Classifier for 80-20 Split

dtree = DecisionTreeClassifier(random_state=13)
dtree.fit(X_train2,y_train2)
y_pred_train2 = dtree.predict(X_train2)
y_pred_test2 = dtree.predict(X_test2)

print('80-20 Model performance on Training Set : \n')
print(classification_report(y_train2,y_pred_train2))
print()
print()
print('80-20 Model performance on Test Set : \n')
```

```

print(classification_report(y_test2,y_pred_test2))
print()
print()
X_ip = list(map(int,
    input("Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : ")
    .split()))[:5]
print('Predicted class : ',dtree.predict([X_ip])[0])
print()
print()
print('Result : ')
pd.DataFrame({'Actual':y_test2,'Predicted':y_pred_test2})

```

80-20 Model performance on Training Set :

	precision	recall	f1-score	support
0	0.97	0.99	0.98	133
1	0.96	0.85	0.90	27
accuracy			0.97	160
macro avg	0.96	0.92	0.94	160
weighted avg	0.97	0.97	0.97	160

80-20 Model performance on Test Set :

	precision	recall	f1-score	support
0	0.74	0.85	0.79	27
1	0.60	0.43	0.50	14
accuracy			0.71	41
macro avg	0.67	0.64	0.65	41
weighted avg	0.69	0.71	0.69	41

Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : 40 0 1 1 0
Predicted class : 0

Result :

Out[13]:

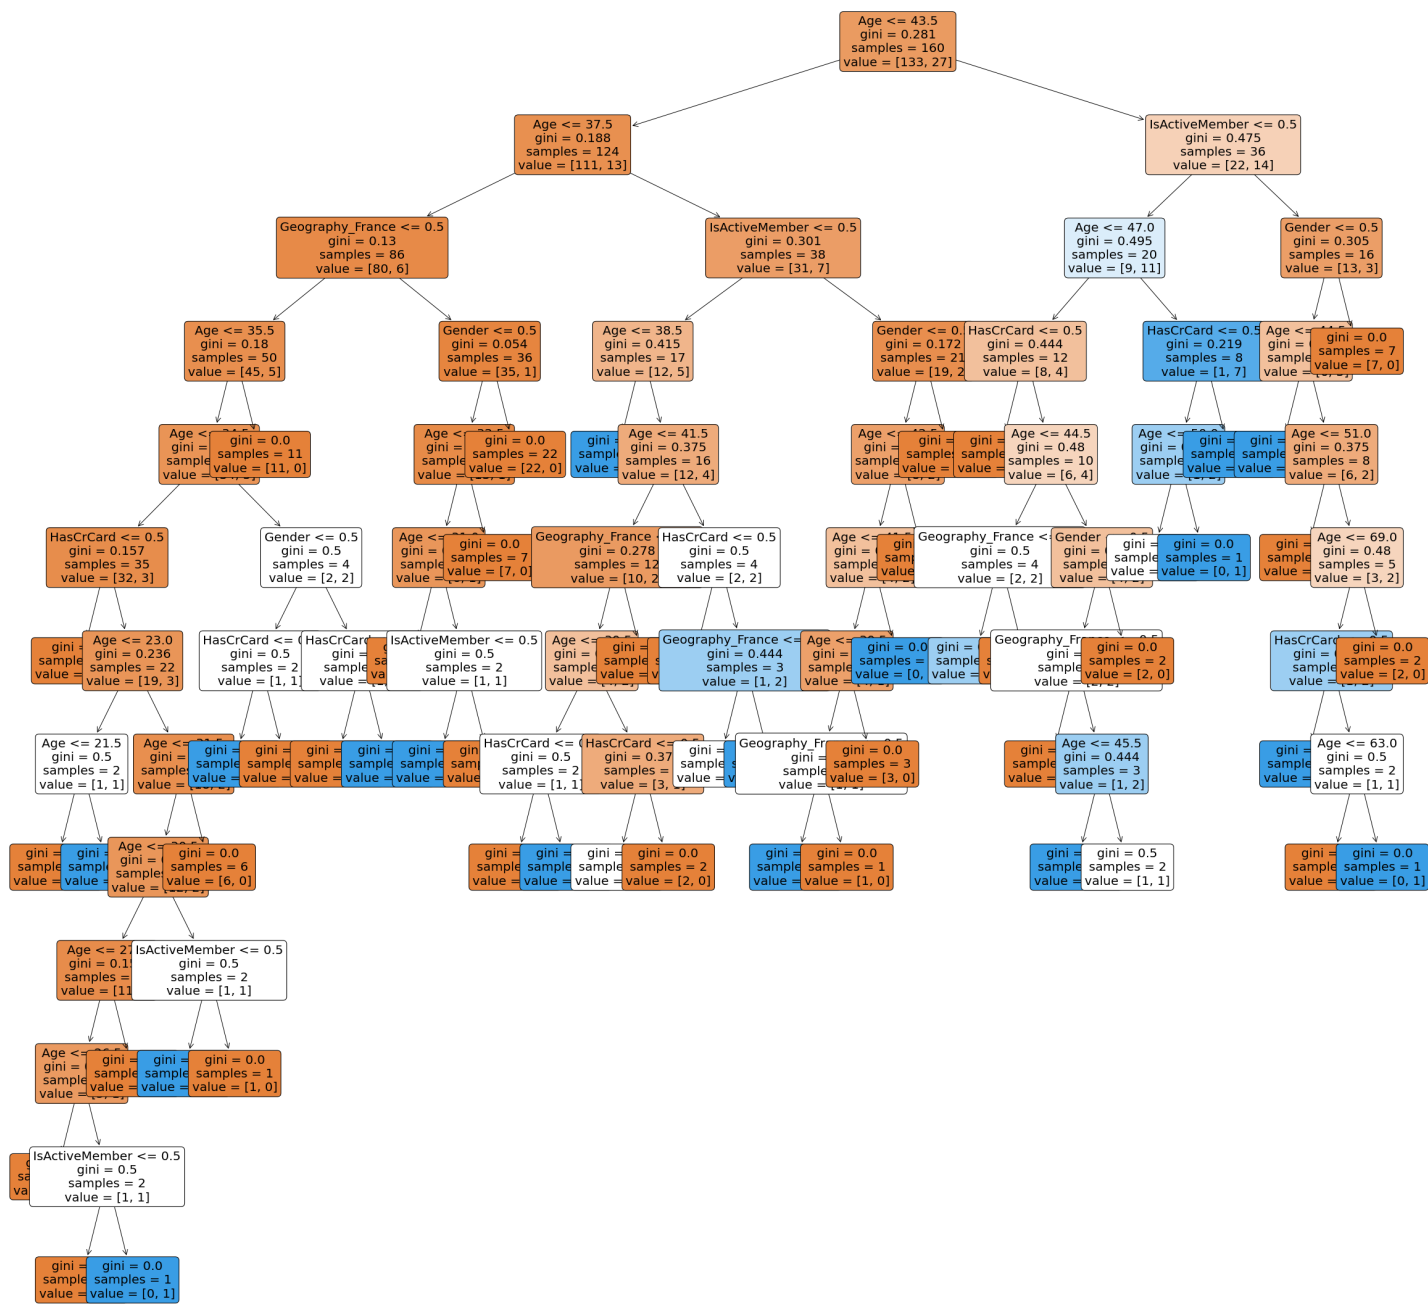
	Actual	Predicted
101	0	0
87	0	0
23	0	0
140	0	1
114	1	0
190	1	0
65	0	0
111	0	1
33	0	1
196	0	0
194	0	0
30	1	1

	Actual	Predicted
70	1	1
169	0	0
164	1	0
132	0	0
103	0	0
179	1	1
105	1	0
184	1	0
191	0	0
143	1	1
68	0	1
7	1	0
125	1	1
175	0	0
124	0	0
144	1	0
127	1	1
166	1	0
79	0	0
97	0	0
102	0	0
106	0	0
10	0	0
92	0	0
13	0	0
42	0	0
15	0	0
131	0	0
95	0	0

In [14]:

```
# Visualize the Tree

plt.figure(figsize=(40,40))
plot_tree(dtree, feature_names = X_train2.columns, filled=True, fontsize=20,
          rounded = True)
plt.show()
```



Training set gave 97% accuracy and test set gave 71% accuracy hence model is slightly overfitted

Decision Tree Classifier (70-30 Split)

In [15]:

```
# Decision Tree Classifier for 70-30 Split

dtree = DecisionTreeClassifier(random_state=13)
dtree.fit(X_train3,y_train3)
y_pred_train3 = dtree.predict(X_train3)
y_pred_test3 = dtree.predict(X_test3)

print('70-30 Model performance on Training Set : \n')
print(classification_report(y_train3,y_pred_train3))
print()
print()
print('70-30 Model performance on Test Set : \n')
print(classification_report(y_test3,y_pred_test3))
print()
print()
X_ip = list(map(int,
    input("Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : ")
```

```

        .split()))[:5]
print('Predicted class : ',dtree.predict([X_ip])[0])
print()
print()
print('Result : ')
pd.DataFrame({'Actual':y_test3,'Predicted':y_pred_test3})

```

70-30 Model performance on Training Set :

	precision	recall	f1-score	support
0	0.97	0.99	0.98	114
1	0.96	0.88	0.92	26
accuracy			0.97	140
macro avg	0.97	0.94	0.95	140
weighted avg	0.97	0.97	0.97	140

70-30 Model performance on Test Set :

	precision	recall	f1-score	support
0	0.83	0.83	0.83	46
1	0.47	0.47	0.47	15
accuracy			0.74	61
macro avg	0.65	0.65	0.65	61
weighted avg	0.74	0.74	0.74	61

Enter Age, isActiveMember, HasCrCard, Gender, Geography_France : 40 0 1 1 0
Predicted class : 0

Result :

Out[15]:

	Actual	Predicted
101	0	0
87	0	0
23	0	0
140	0	1
114	1	0
...
32	0	1
19	0	0
17	0	0
173	0	0
193	0	0

61 rows × 2 columns

In [16]: `# Visualize the Tree`

